

Lesson 11 - zkEVM Solutions

zkEVM Solutions

Rollup Recap

Rollups are solutions that have

- transaction execution outside layer 1
- transaction data and proof of transactions is on layer 1
- a rollup smart contract in layer 1 that can enforce correct transaction execution on layer 2 by using the transaction data on layer 1

The main chain holds funds and commitments to the side chains

The side chain holds additional state and performs execution

There needs to be some proof, either a fraud proof (Optimistic) or a validity proof (zk)

Rollups require “operators” to stake a bond in the rollup contract. This incentivises operators to verify and execute transactions correctly.

Data Availability in general

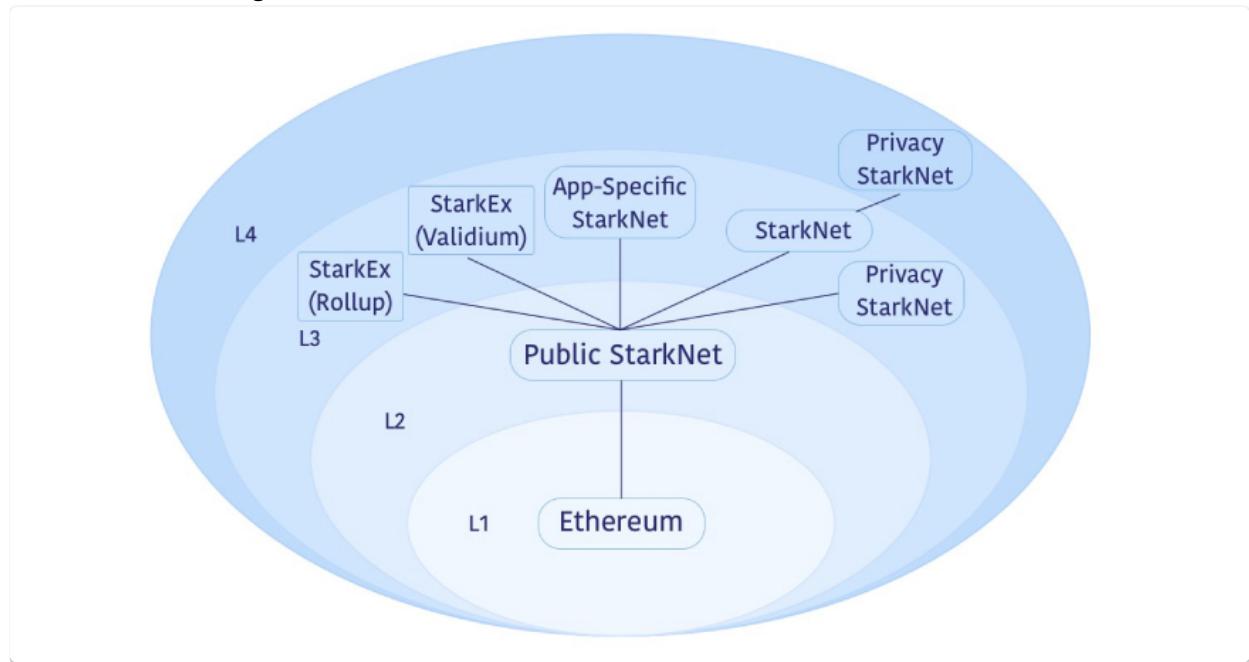
In order to re-create the state, transaction data is needed, the data availability question is where this data is stored and how to make sure it is available to the participants in the system.

	Validity Proofs		Fault Proofs
Data On-Chain	Volition	ZK-Rollup	Optimistic Rollup
Data Off-Chain		Validium	Plasma

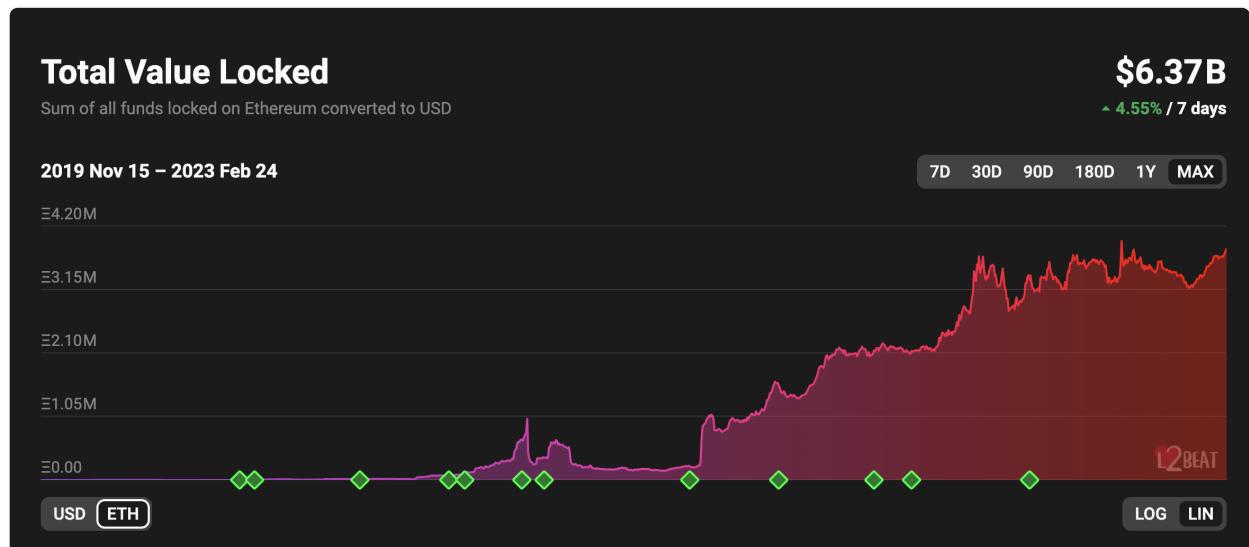
STARKWARE

L3 and L4 ?

See Fractal scaling article

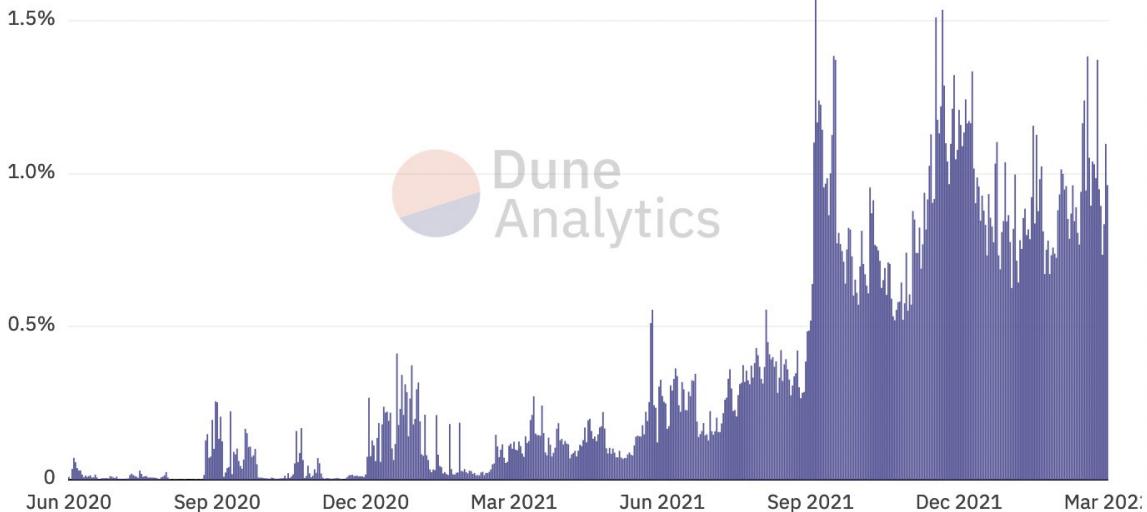


L2 Statistics



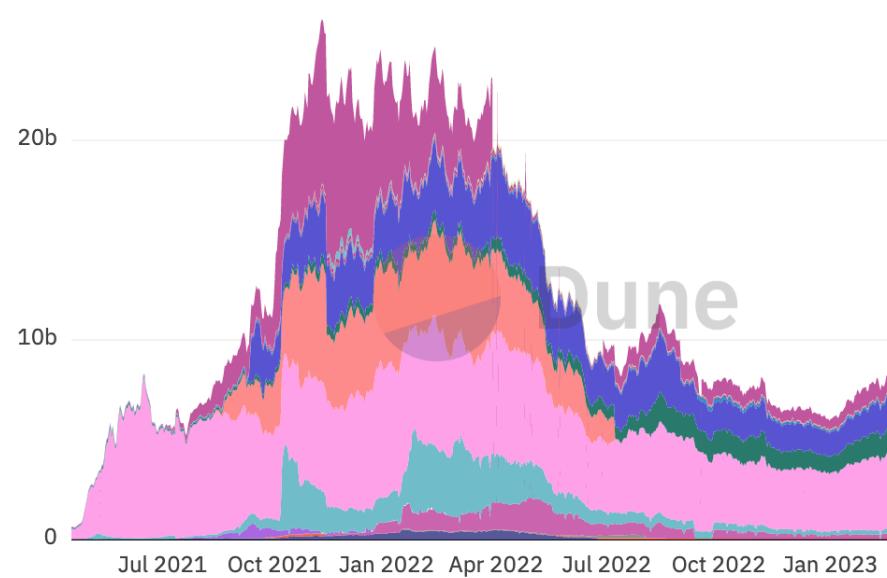
% gas spent by l2 proof contracts of the total max daily available ethereum gas
Actual ethereum daily limit is : 100,000,000,000 gas

@funnyking



Ethereum bridges TVL over time ↗

@eliasimos



- Harmony Bridges
- Synapse Bridge
- Nomad Bridge
- Near Rainbow Bridge
- Solana Wormhole
- Fantom Anyswap
- Polygon Bridges
- Avalanche Bridge
- Optimism Bridges
- Arbitrum Bridges
- xDAI Bridges
- Optics Bridge
- BSC Anyswap Brid
- Moonriver Anyswa
- RSK Token Bridge
- Boba Network Brid

TVL DAI USDT



Transaction compression and costs

For zk rollups it is expensive to verify the validity proof.

For STARKs, this costs ~5 million gas, which when aggregated is about 384 gas cost per transaction.

Due to compression techniques, the calldata cost is actually only 86 gas.

This is further reduced by the calldata [EIP448](#) to 16.1 gas.

The batch cost is poly-log, so if activity increases 100x, the batch costs will decrease to only 4–5 gas per transaction, in which case the calldata reduction would have a huge impact.

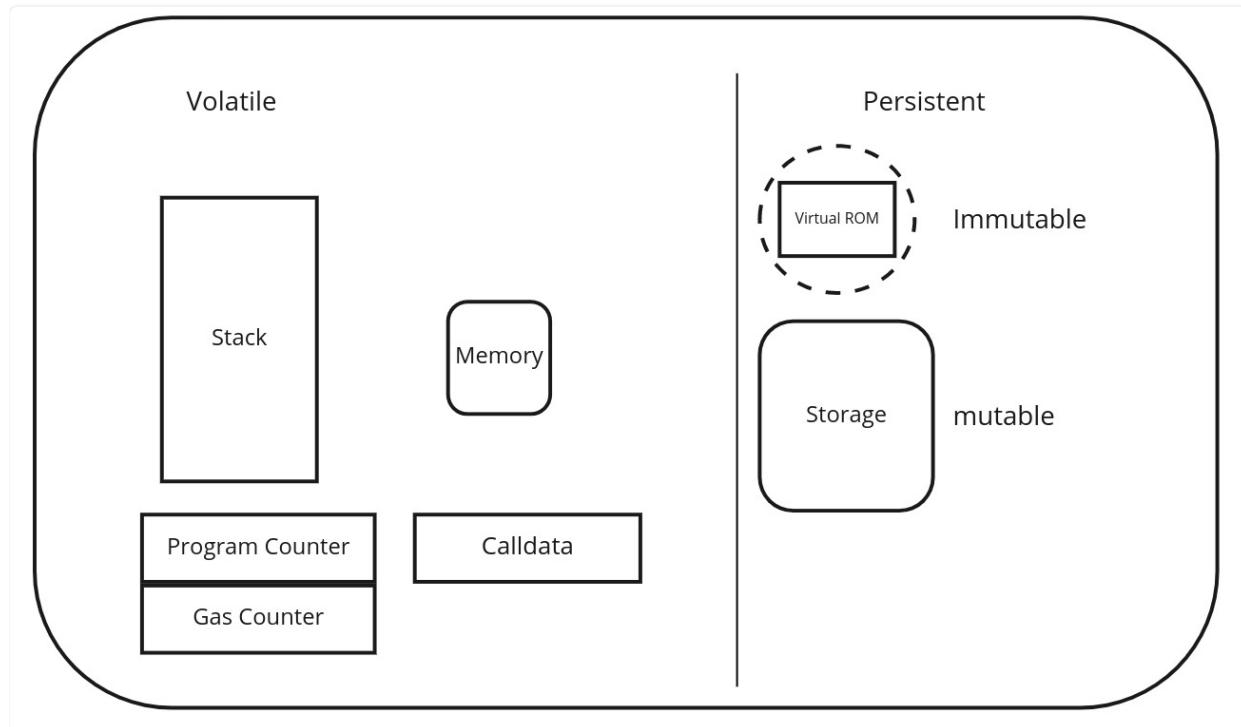
At 100x the TPS today on dYdX, the total on-chain cost will reduce to only 21 gas

At this point, the bottleneck becomes prover costs for the rollup as much as on-chain gas fees, see [Polygon Zero](#) and recursive proofs

Approaches to zkRollups on Ethereum

1. Building application-specific circuit (although this can be fairly generic as in Starknet)
 2. Building a universal “EVM” circuit for smart contract execution
-

zkEVM Solutions in general



The opcode of the EVM needs to interact with Stack, Memory, and Storage during execution. There should also be some contexts, such as gas/program counter, etc. Stack is only used for Stack access, and Memory and Storage can be accessed randomly.

EVM processing

In general the EVM will

1. Read elements from stack, memory or storage
2. Perform some computation on those elements
3. Write back results to stack, memory or storage

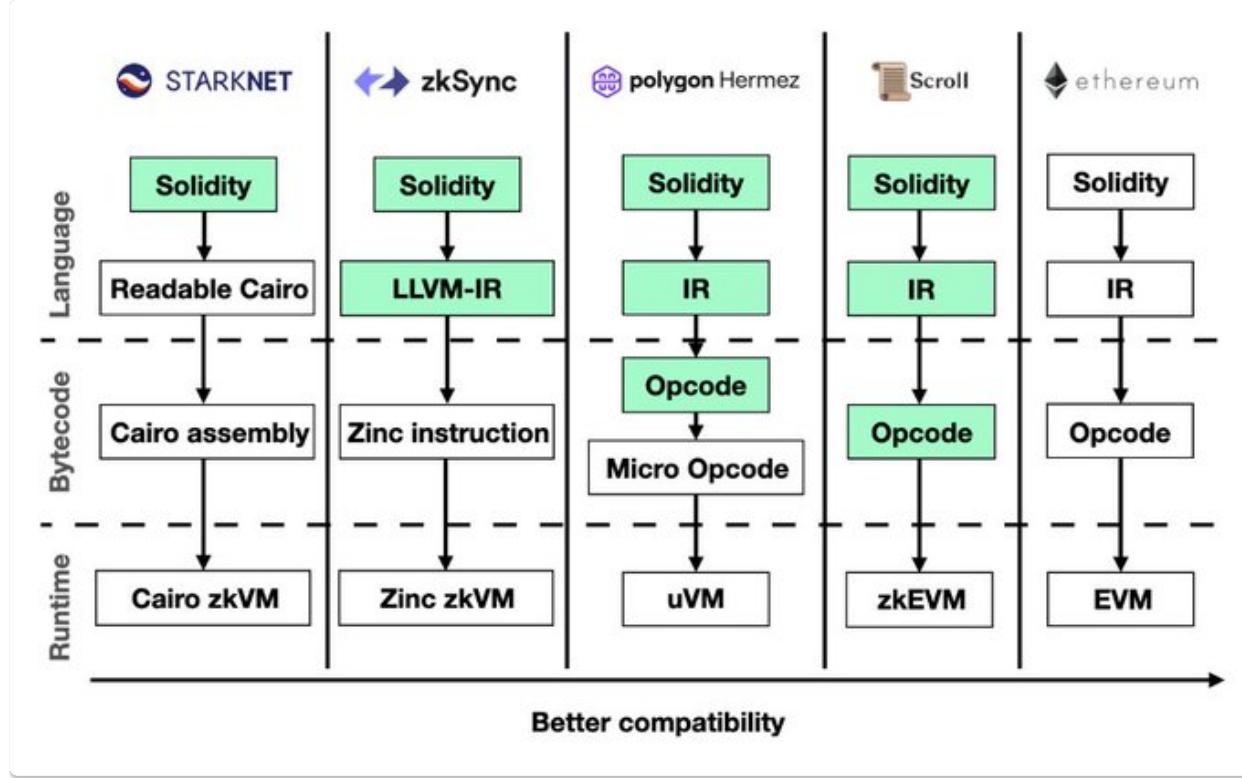
So our circuit has to model / prove this process, in particular

- The bytecode is correctly loaded from persistent storage
- The opcodes in the bytecode are executed in sequence
- Each opcode is executed correctly (following the above 3 steps)

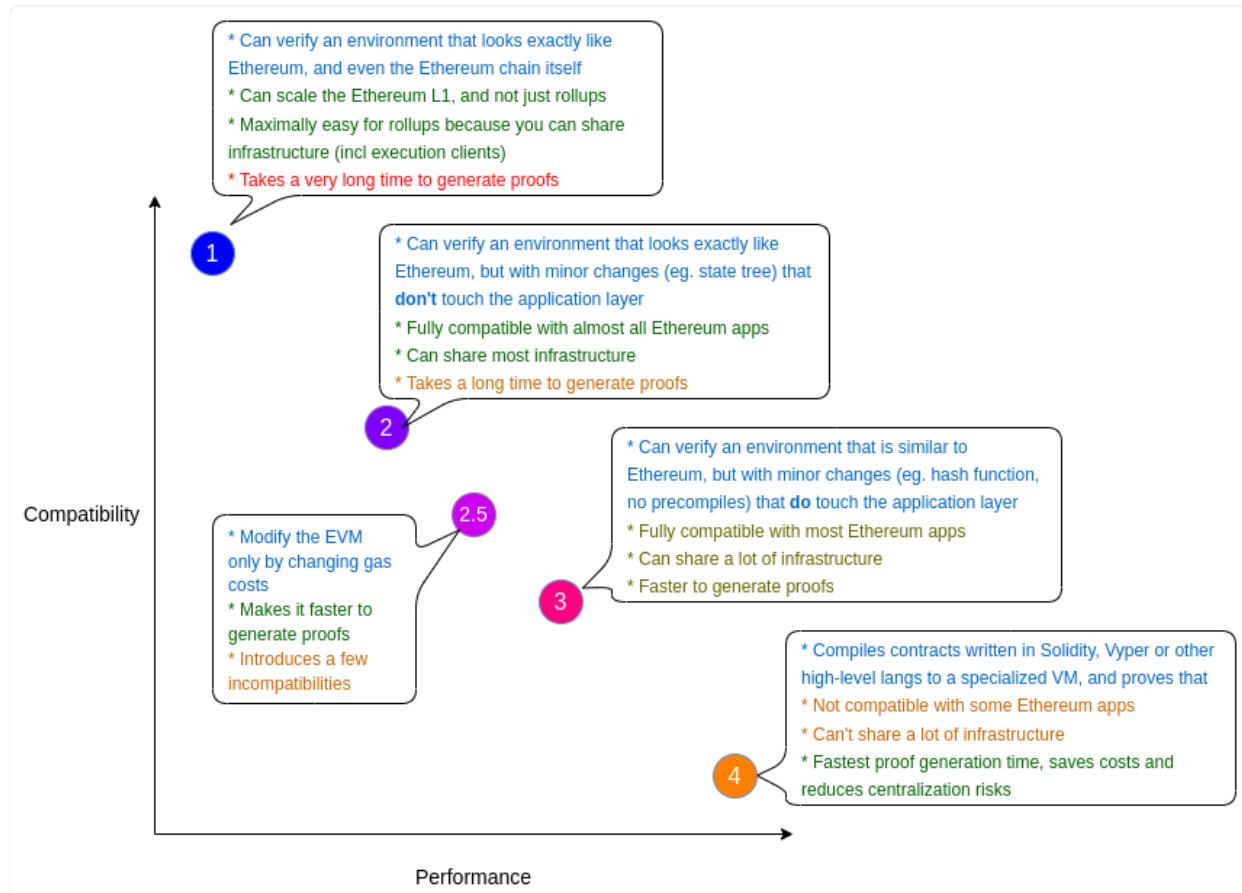
Design challenges in designing a zkEVM

1. We are constrained by the cryptography (curves, hash functions) available on Ethereum.
2. The EVM is stack based rather than register based
3. The EVM has a 256 bit word (not a natural field element size)
4. EVM storage uses keccak and Merkle Patricia trees, which are not zkp friendly
5. We need to model the whole EVM to do a simple op code.

zkEVM taxonomy



See Vitalik article



Type 1 (fully Ethereum-equivalent)

See zkEVM research [team](#)

Type 2 (fully EVM-equivalent)

(not quite Ethereum-equivalent)

[Scroll](#)

[Hermez](#)

Type 2.5 (EVM-equivalent, except for gas costs)

Type 3 (almost EVM-equivalent)

Scroll and Hermez in their current form

Type 4 (high-level-language equivalent)

[ZKSync](#)

[Starknet + Warp](#)

Also see

The [ZK-EVM Community Edition](#) (bootstrapped by community contributors including [Privacy and Scaling Explorations](#), the Scroll team, [Taiko](#) and others) is a Tier 1 ZK-EVM.

zkSync

Resources

[Web](#)

[Twitter](#)

[Medium](#)

[Docs](#)



ABILITIES

3 unique

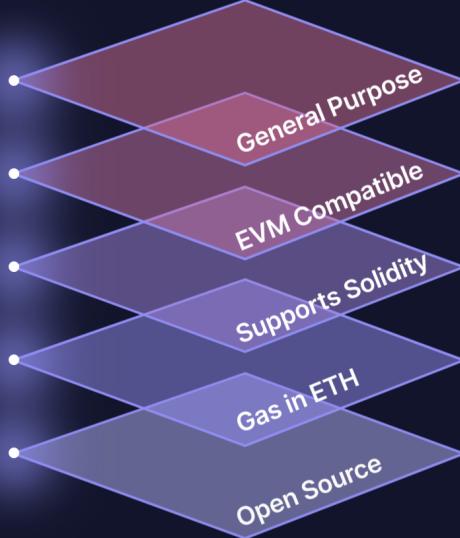
zkSync 2.0 supports Account Abstraction, Solidity and Vyper, and is built with an LLVM compiler that will one day give us the ability to tap into the power of other programming libraries written in Rust, C++, and Swift. In addition, any project built on our Layer 2 solution will be future-proof when we launch our Layer 3 solution.

[Build](#) [Explore](#)

INGREDIENTS

5 magical

We've all been waiting a long time for a zkRollup that represents the end-game for Ethereum scaling - one that scales the technology and values of Ethereum without degrading security or decentralization. It's taken 4 long years, but **zkSync 2.0** is now on mainnet, rapidly moving toward a full gated-release.



[Build](#) [Explore](#)

Overview

zkSync is built on ZK Rollup architecture. ZK Rollup is an L2 scaling solution in which all funds are held by a smart contract on the mainchain, while computation and storage are performed off-chain.

For every Rollup block, a state transition zero-knowledge proof is generated and verified by the mainchain contract.

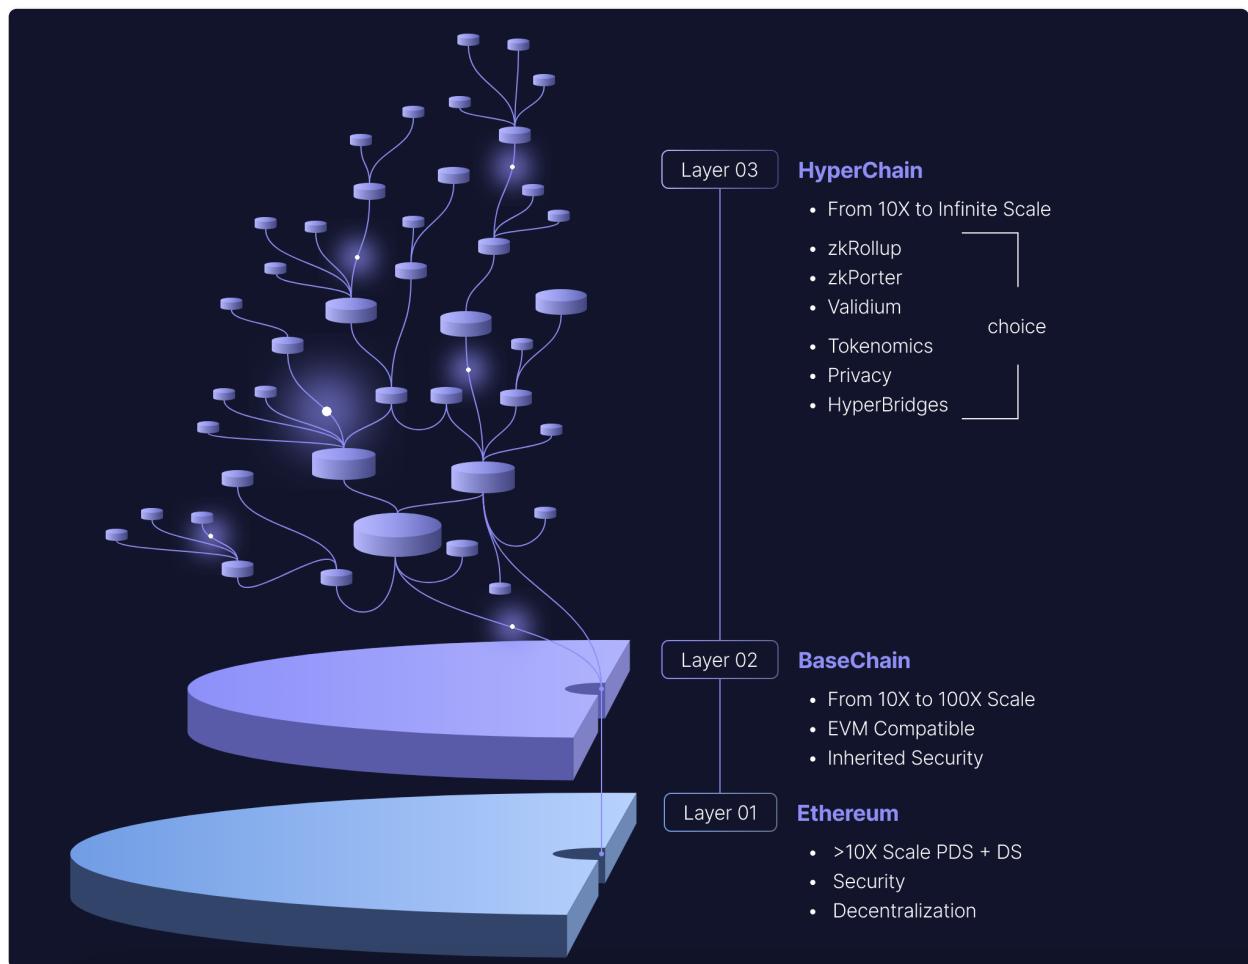
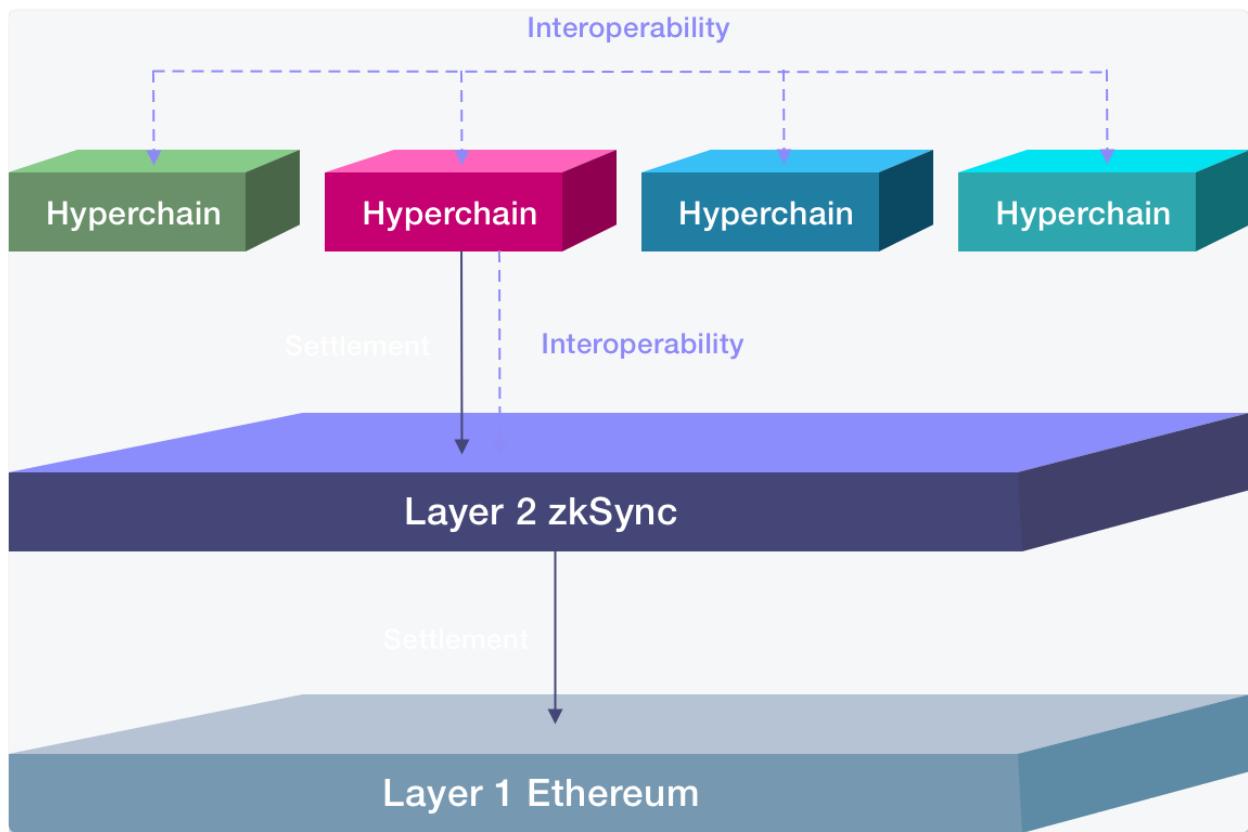
This SNARK includes the proof of the validity of every single transaction in the Rollup block. Additionally, the public data update for every block is published over the mainchain network in the cheap calldata.

Hyperscaling

Hyperchains are fractal-like instances of zkEVM running in parallel and with the common settlement on the L1 mainnet.

The Basechain is the main Hyperchain instance of zkSync Era (the L2 instance). It serves as the default computation layer for generic smart contracts and as a settlement layer for all other Hyperchains (L3 and above).

The Basechain is not special in any particular way except that it settles its blocks directly on the L1



Sequencing transactions

Different modes are available

Centralized sequencer - In this mode, there will be a single centralized operator with a conventional REST API to accept transactions from users.

Decentralized sequencer - In this mode, a Hyperchain will coordinate on what transactions are included in a block using a consensus algorithm.

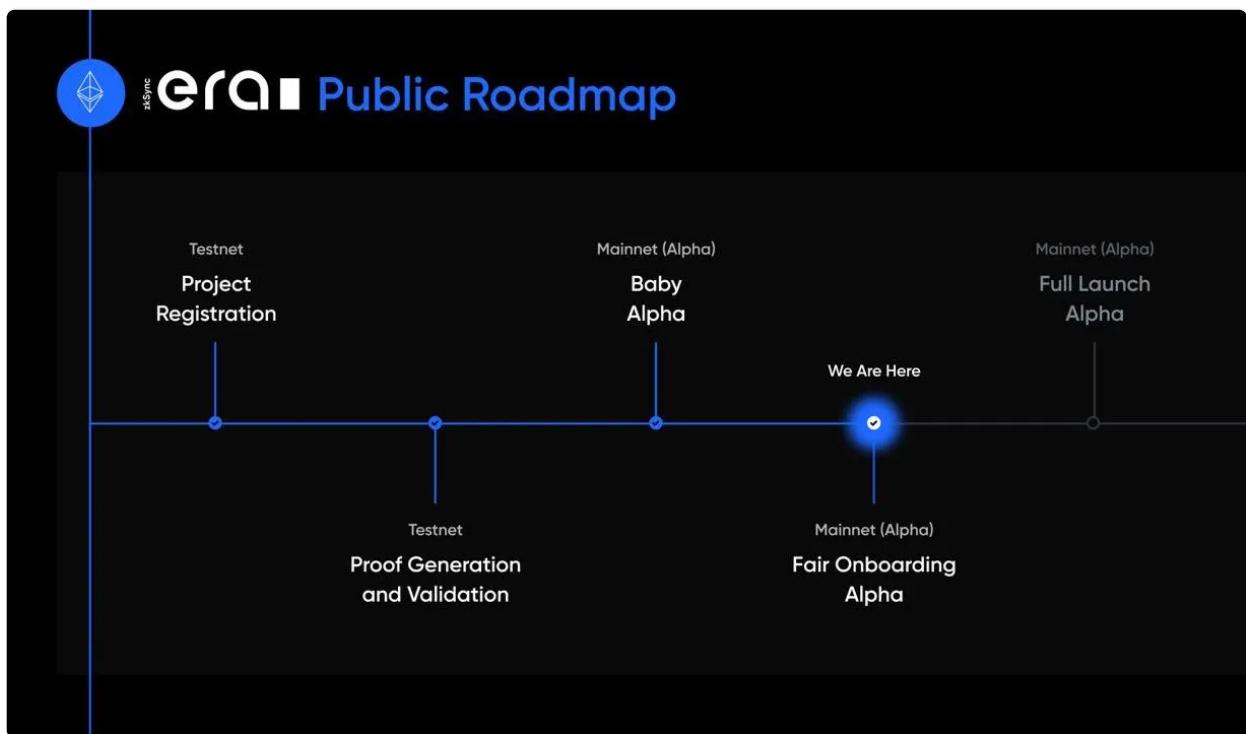
Any algorithm is possible, such as Tendermint or HotStuff with permissionless dPoS.

Priority queue - The goal of the priority queue is to provide a censorship-resistant way to interact with zkSync in case the operator becomes malicious or unavailable. all transactions can be submitted in batches via the priority queue from an underlying L2 or even L1 chain.

zkSync Roadmap

zkSync 2.0 is now zkSync Era

zkSync era



From blog post

We've arrived at **Fair Onboarding Alpha**, where we welcome projects that have registered to launch on zkSync Era. Fair onboarding is designed to be equitable, so we are welcoming registered projects at the same time, giving all teams an equal opportunity to deploy.

During Fair Onboarding Alpha, zkSync will

- allow registered projects to deploy and test their dapps on zkSync Era
- enable token bridging, with limited use for testing and deployment purposes
- continue running security audits, contests, and bug bounty programs
- improve developer tools, including plugins, documentation, tutorials and FAQs.

Data availability options on zkSync

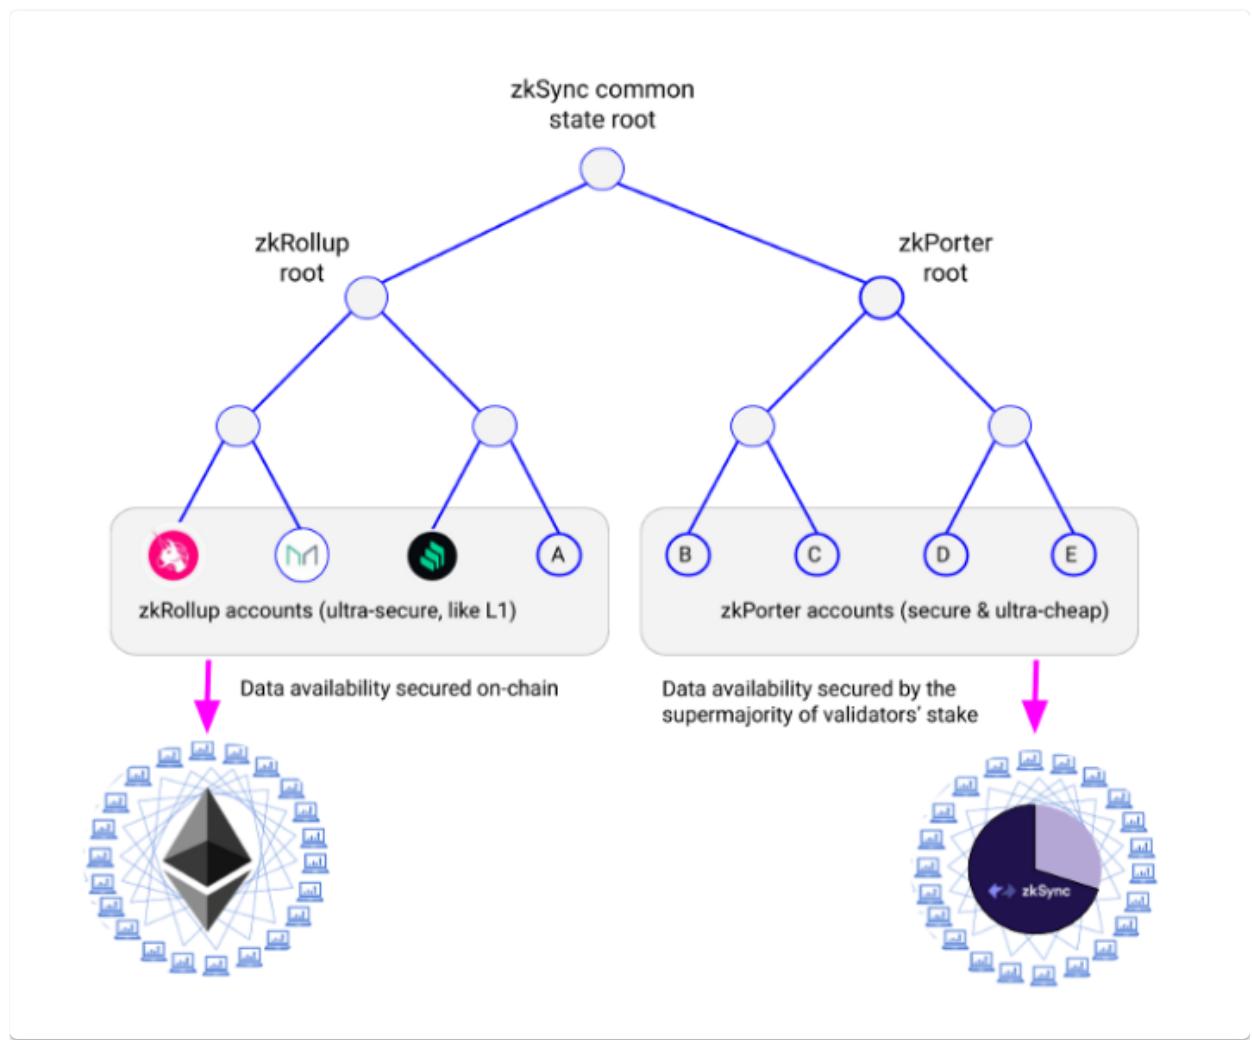
- zkRollup - The default policy

The values of every changed storage slot at the end of the block must be published as calldata on L1.

Note that it is the end state change that is published, if there were multiple transactions contributing to the change, then the cost can be amortised.

zkPorter - off chain data availability, see this [article](#)

contracts and accounts on the zkRollup side will be able to seamlessly interact with accounts on the zkPorter side.

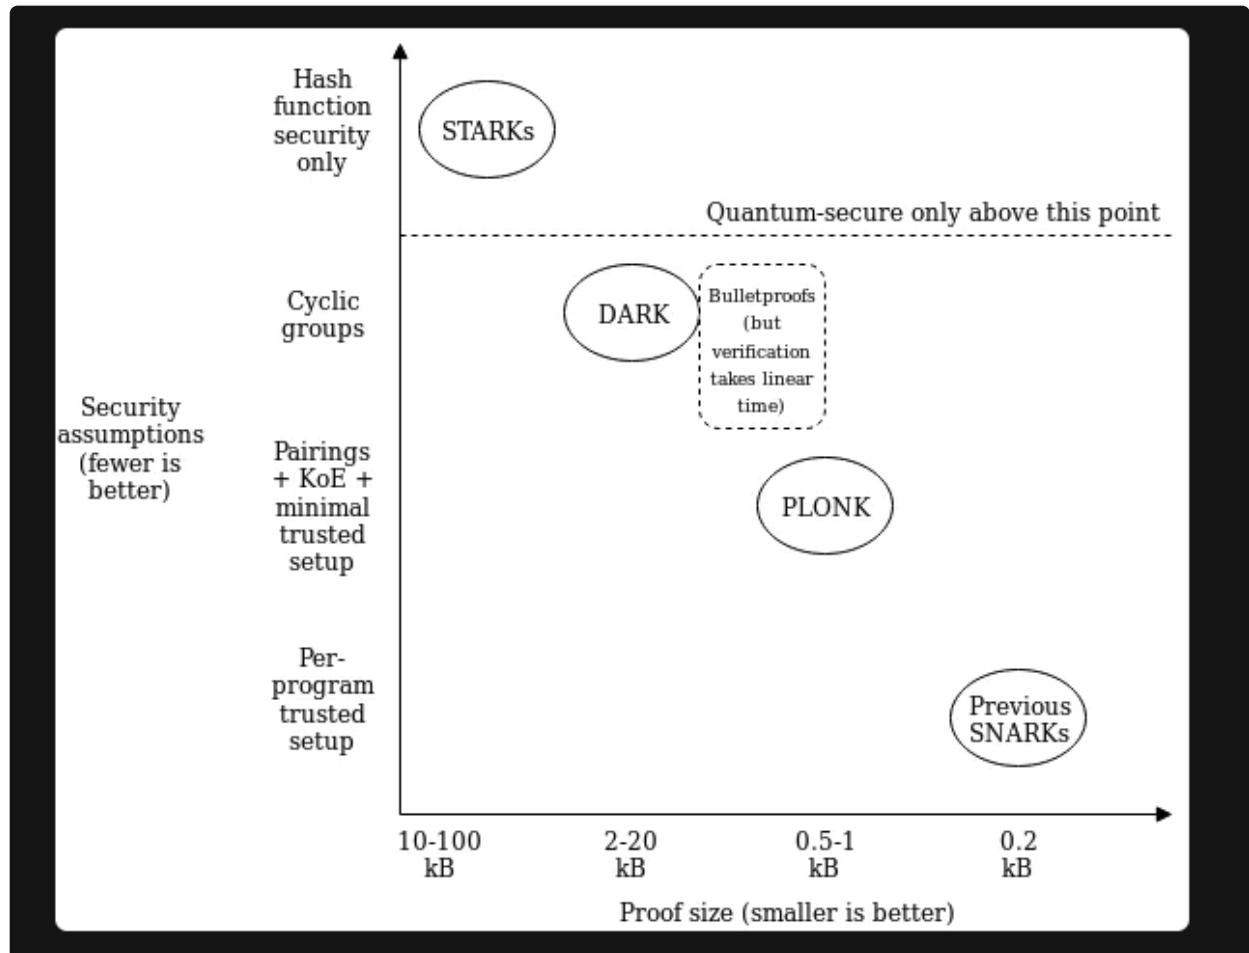


From the [article](#)

"Uniswap deploys their smart contract on the zkRollup side, and retail users on a zkPorter account can swap for <\$0.03 in fees.

The overwhelming majority of rollup fees are due to the costs of publishing data on Ethereum. zkPorter accounts can make thousands of swaps on the Uniswap contract, but only a single update needs to be published to Ethereum."

Underlying Cryptography



zkSync uses PLONK with custom gates and lookup tables (most commonly referred to as UltraPLONK) and Ethereum's BN-254 curve.

zkSync Infrastructure

zkSync operates several pieces of infrastructure on top of Ethereum. All infrastructure is currently live and operational, including the zkEVM.

- Full Node
 - Executes zkEVM bytecode using the virtual machine
 - Filters incorrect transactions
 - Executes mempool transactions
 - Builds blocks
- Prover
 - Generates ZK proofs from block witnesses
 - provides an interface for parallel proof generation
 - Scalable (can increase # of provers depending on demand)
- Interactor
 - The link between L1 Ethereum and L2 zkSync
 - Calculates transaction fees

- Fees depend on token prices, proof generation, and L1 gas costs
- Paranoid Monitor
 - Monitors infrastructure and notifies Matter Labs if incidents occur

zkSync Ecosystem

<https://ecosystem.zksync.io/>

Faucet for test net

New Block explorer

Implementation code on L1 at 0xd61dFf4b146e8e6bDCDad5C48e72D0bA85D94DbC

Block proving contract

```
/// @notice Blocks commitment verification.
/// @notice Only verifies block commitments without any other processing
function proveBlocks(StoredBlockInfo[] memory _committedBlocks,
ProofInput memory _proof) external nonReentrant {

    requireActive();
    uint32 currentTotalBlocksProven = totalBlocksProven;
    for (uint256 i = 0; i < _committedBlocks.length; ++i) {

        require(hashStoredBlockInfo(_committedBlocks[i]) ==
storedBlockHashes[currentTotalBlocksProven + 1], "o1");

        ++currentTotalBlocksProven;
        require(_proof.commitments[i] & INPUT_MASK ==
uint256(_committedBlocks[i].commitment) & INPUT_MASK, "o"); // incorrect
block commitment in proof

    }
    bool success =
verifier.verifyAggregatedBlockProof(
    _proof.recursiveInput,
    _proof.proof,
    _proof.vkIndexes,
    _proof.commitments,
    _proof.subproofsLimbs
);

    require(success, "p"); // Aggregated proof verification fail
    require(currentTotalBlocksProven <= totalBlocksCommitted, "q");
    totalBlocksProven = currentTotalBlocksProven;

}
```

Comparing ZkSync with StarkEx

zkSync vs. StarkWare

	zkSync	StarkWare
Proof	zk-Rollups	zk-Rollups or Validium
Optimal throughput	~300 or 800 - 3000	~3000
Txs/Proof	100 or 315	300
Prover time (minutes)	4 - 14	3 - 5
Optimal gas cost	~1200	~300
Fixed-cost in gas	500k - 900k	~2.5M-5M
Txs weekly on ETH	44k	3.95M
TVL	22M	1B
Hardware	Customized FPGA	Own prover device
Data availability	Yes	Yes or Committee
Software license	Open-source (Apache/MIT)	Not open-source (others cannot run STARK prover)
Developer stack	Focus on EVM	Built Cairo (not backwards compatible)
Upgradeability	Mandatory Timelock	Freeze operation until upgrade executed

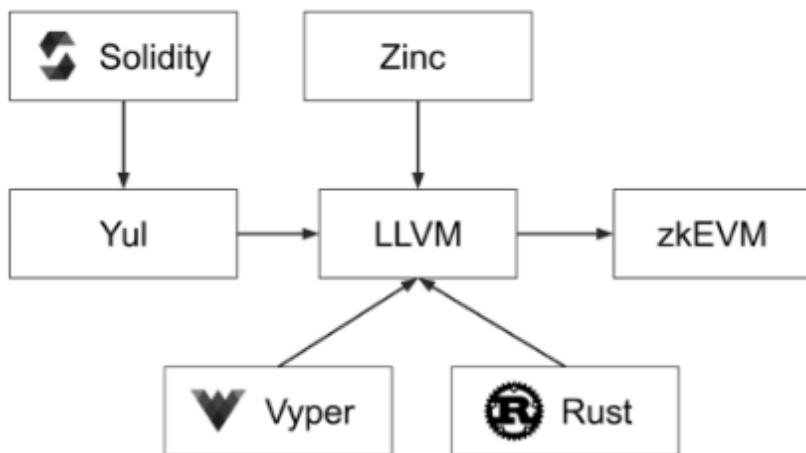
Matter Labs zkEVM

zkEVM is a virtual machine that executes smart contracts in a way that is compatible with zero-knowledge-proof computation.

zk-EVM keeps EVM semantics, but is also ZK-friendly and takes on traditional CPU architectures.

zkSync's zkEVM is not a replica of the EVM but is newly designed to run 99% of Solidity contracts and ensure that it works properly under a variety of conditions (including rollbacks and exceptions). At the same time, zkEVM can be used to efficiently generate zero-knowledge proofs in the circuit.

zkEVM compiler



The circuit implementation of Matter Labs uses TinyRAM to implement ordinary opcodes, such as ADD, PUSH, etc.; opcodes that consume a lot of gas, such as SHA256/keccak, implement this circuit especially; finally, Matter Labs uses recursive aggregation technology to aggregate all proofs into one proof.

Polygon Products

See this [guide](#)

Strategy [article](#)

Recent [paper](#) on efficient zk proofs for Keccak

Polygon Zero

zkRollup solutions have a bottleneck in the time it takes to generate a proof.

Polygon Zero attempts to solve this with "recursive proofs", based on [Plonky2](#).

Polygon Zero generates proofs simultaneously for every transaction in the batch. These are then aggregated into a single proof which is submitted on the Ethereum network.

This approach significantly reduces the effort it takes to generate reliable validity proofs. Polygon Zero's Plonky2 can generate a recursive proof in 0.17 seconds.

Impact Statistics

Implementing Polygon Zero results in below numbers.

0.17 sec

To generate a zero-knowledge proof

45kb

Size of proofs

~5bits

Amount of data stored by validators per active account

Polygon Hermez

Polygon (Hermez) team are working on a protocol Proof of Efficiency

This involves 2 permissionless roles :

- Sequencer
- Aggregator

From [article](#)

Sequencers collect transactions from users on the rollup, then select and pre-process new batches of this Layer 2 data. Finally, they send transactions to Layer 1 to be recorded. Sequencers also deposit a fee in \$MATIC token as an incentive for Aggregators to include the batch in a zero knowledge proof.

Hermez 2.0

Hermez current functionality is limited to token transfers and atomic swaps, so there are plans to introduce Hermez 2.0 which will have EVM compatibility.

Polygon Miden

Polygon Miden is a general-purpose, STARK-based ZK rollup with EVM compatibility.

This will differ from Starknet in that it will be EVM compatible, so it should run Solidity contracts.

From their [documentation](#)

"Polygon Miden can process up to 5,000 transactions in a single block, with new blocks produced every five seconds. Although this ZK rollup exists as a prototype for now, it is expected to boost throughput to over 1,000 transactions per second (TPS) at launch."

Polygon Nightfall

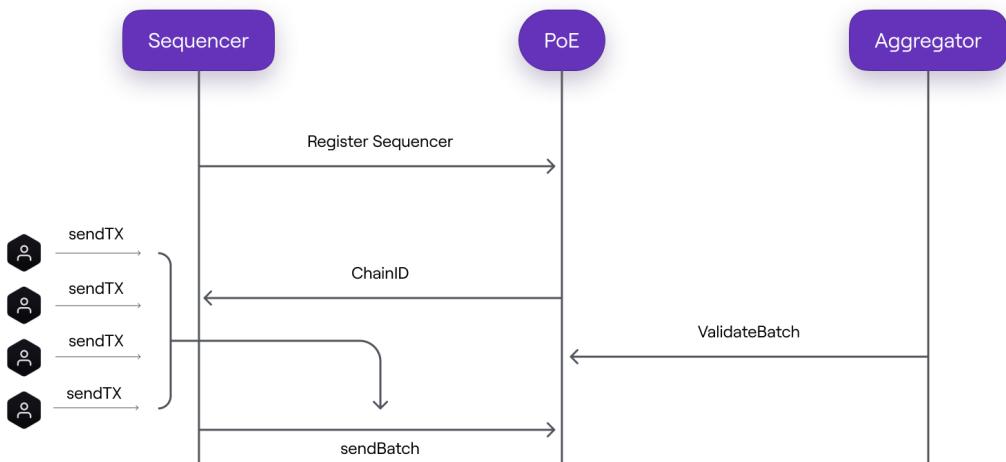
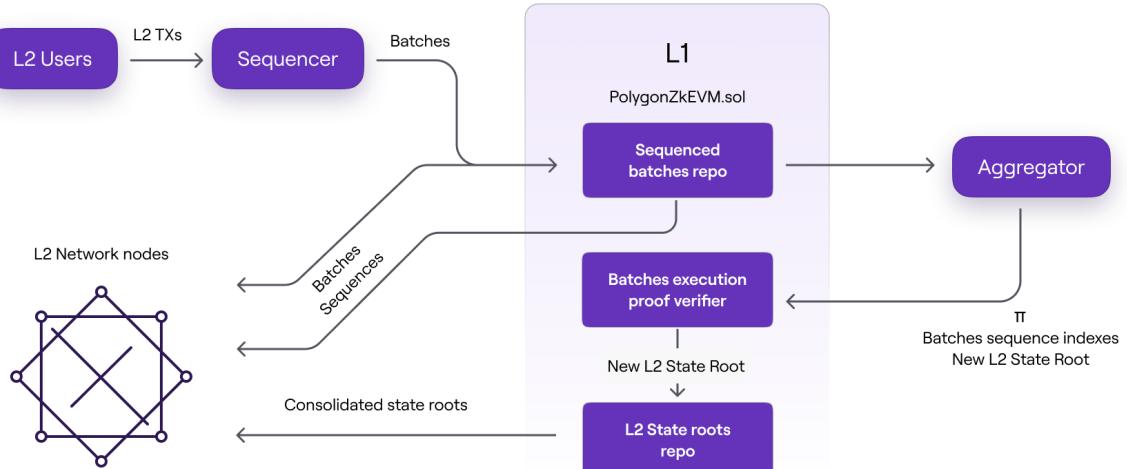
From a collaboration with EY it is designed to allow private transactions.

It is a combination Optimistic rollups and zero knowledge, optimistic rollups for scalability and zk for privacy.

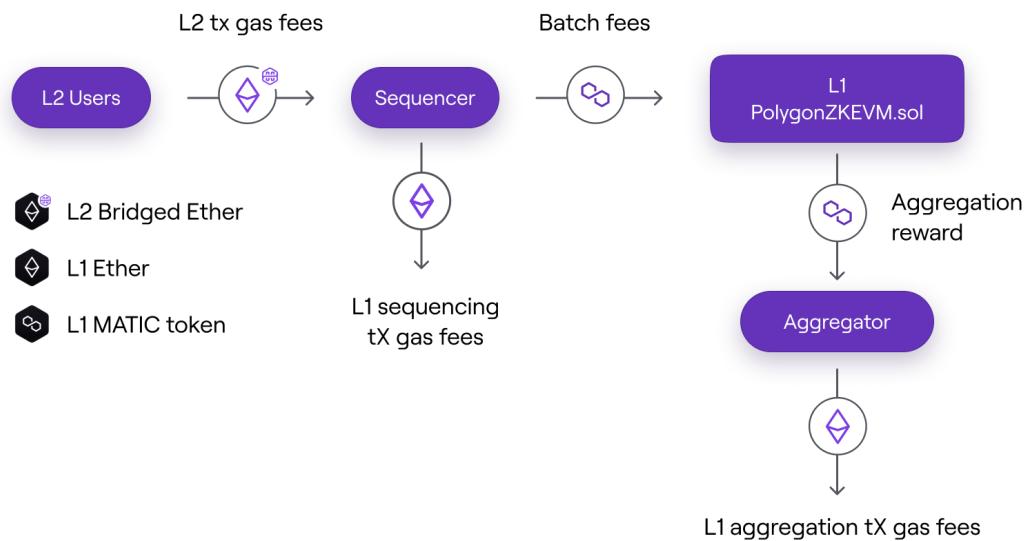
There is a beta version available on mainnet.

Polygon zkEVM

Transaction / State Flow



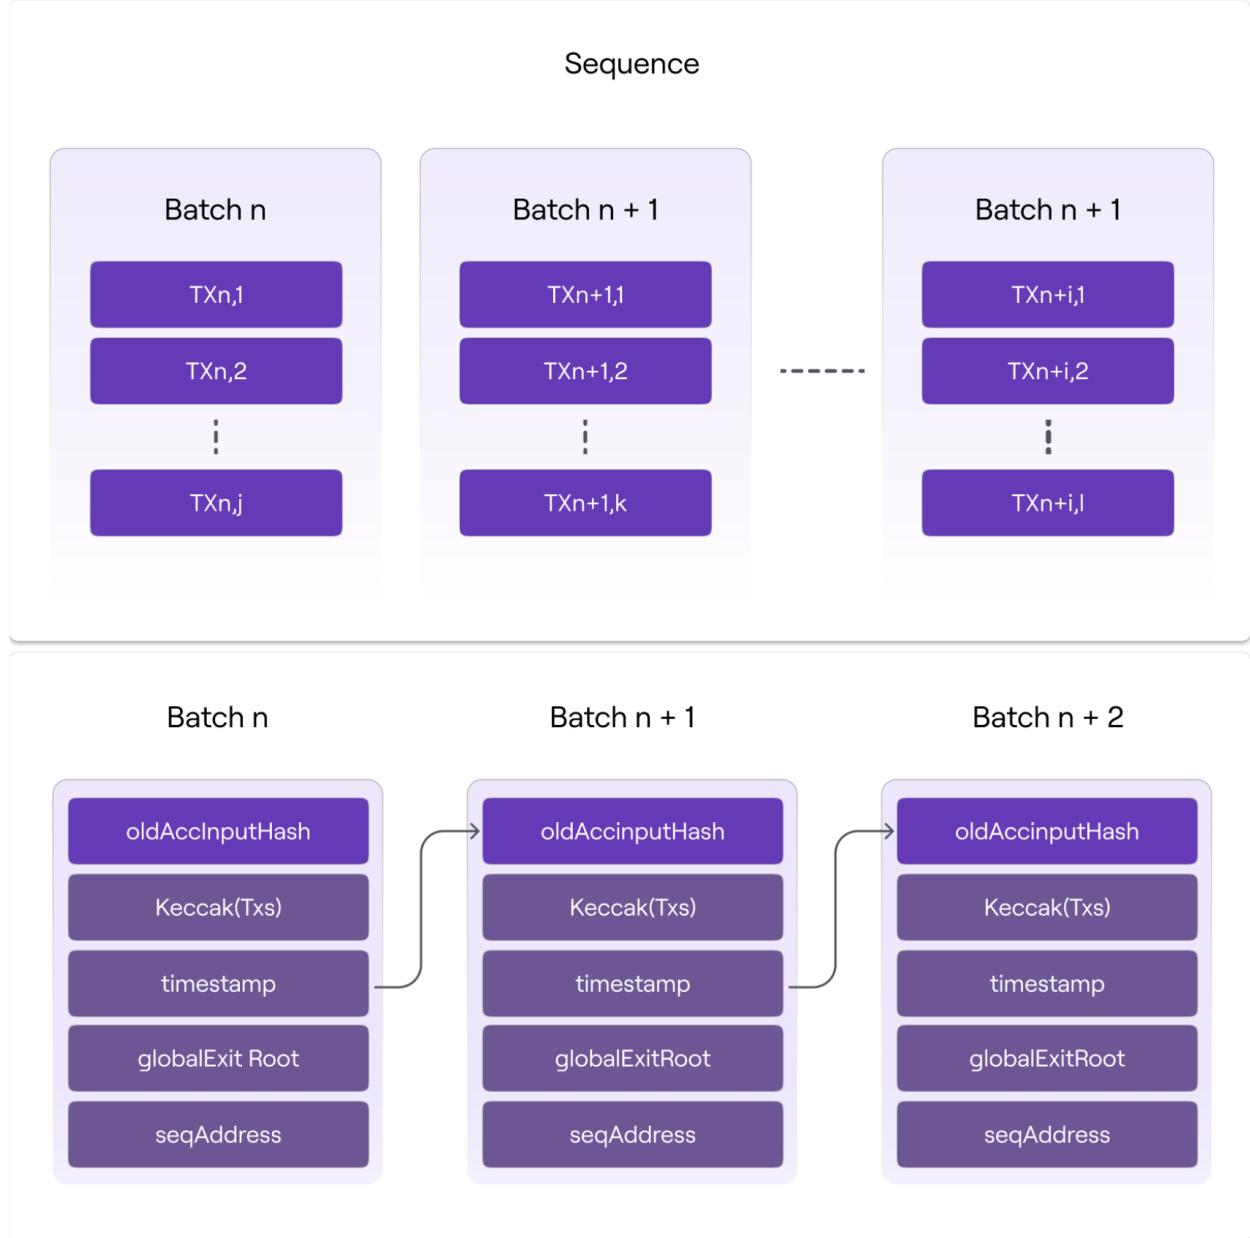
Fees



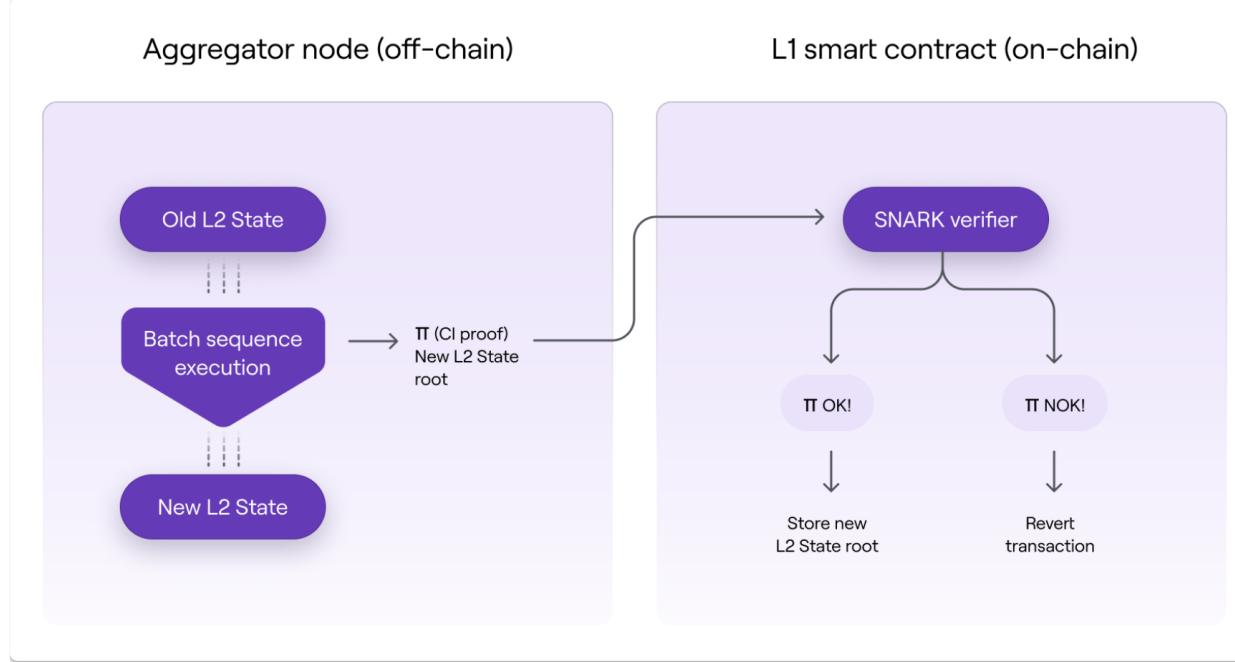
The Sequencer prioritizes transactions with higher gas prices. Furthermore, there is a threshold below which it is unprofitable for the Sequencer to execute transactions because the fees earned from L2 users are less than the fees paid for sequencing fees (plus L1 sequencing transaction fee).

The aggregator also receives an aggregation reward.

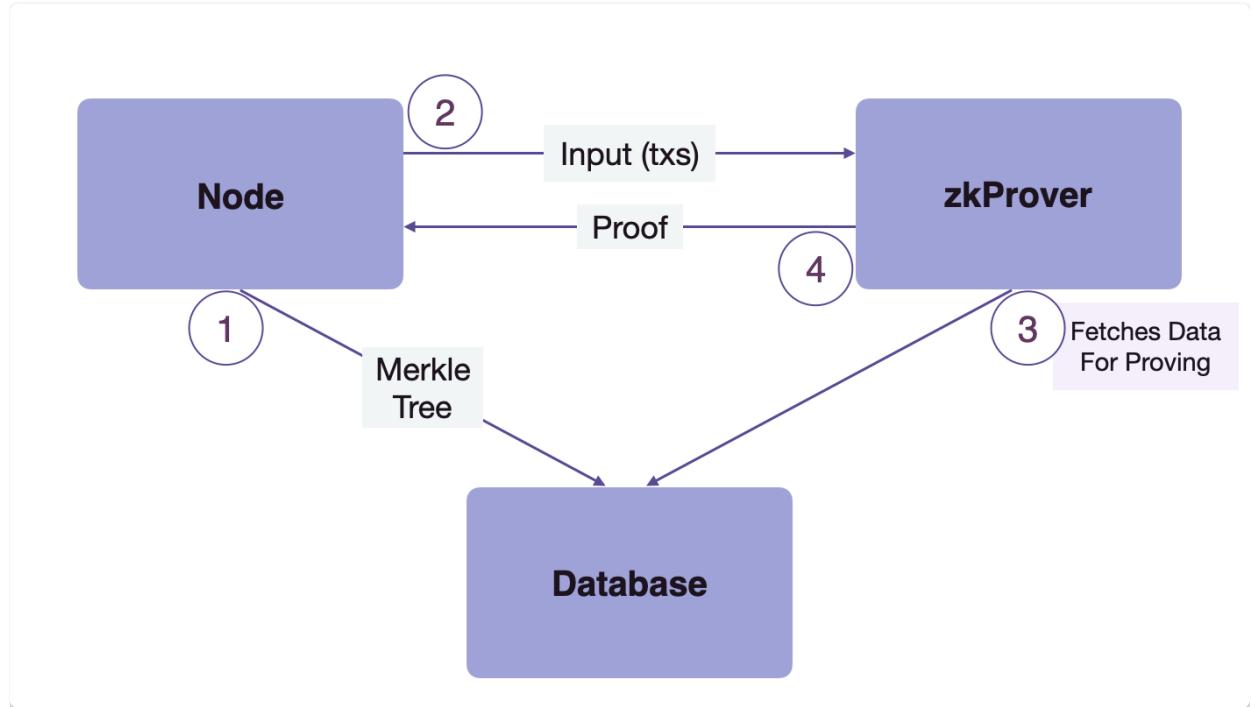
Batch Sequencing



Batch Aggregation



Proof process



In summary the process is

- The Node sends the content of Merkle trees to the Database to be stored there
- The Node then sends the input transactions to the zkProver
- The zkProver accesses the Database and fetches the info needed to produce verifiable proofs of the transactions sent by the Node. This information consists of the Merkle roots, the keys and hashes of relevant siblings, and more
- The zkProver then generates the proofs of transactions, and sends these proofs back to the Node

See [Docs](#)

Polynomial Identity Language

See [Docs](#)

From documentation

This is a novel domain-specific language useful for defining state machines. The aim for creating PIL is to provide developers a holistic framework for both constructing state machines through an easy-to-use interface, and abstracting the complexity of the proving mechanisms.

One of the main peculiarities of PIL is its **modularity**, which allows programmers to define parametrizable state machines, called namespaces, which can be instantiated from larger state machines. Building state machines in a modular manner makes it easier to test, review, audit and formally verify even large and complex state machines. In this regard, by using PIL, developers can create their own custom namespaces or instantiate namespaces from some public library.

Some of the keys features of PIL are;

- Providing namespaces for naming the essential parts that constitute state machines,
- Denoting whether the polynomials are committed or constant,
- Expressing polynomial relations, including identities and lookup arguments and
- Specifying the type of a polynomial, such as bool or u32u32.

Some example PIL



We will look at this further in our lesson about PLONK.

Scroll

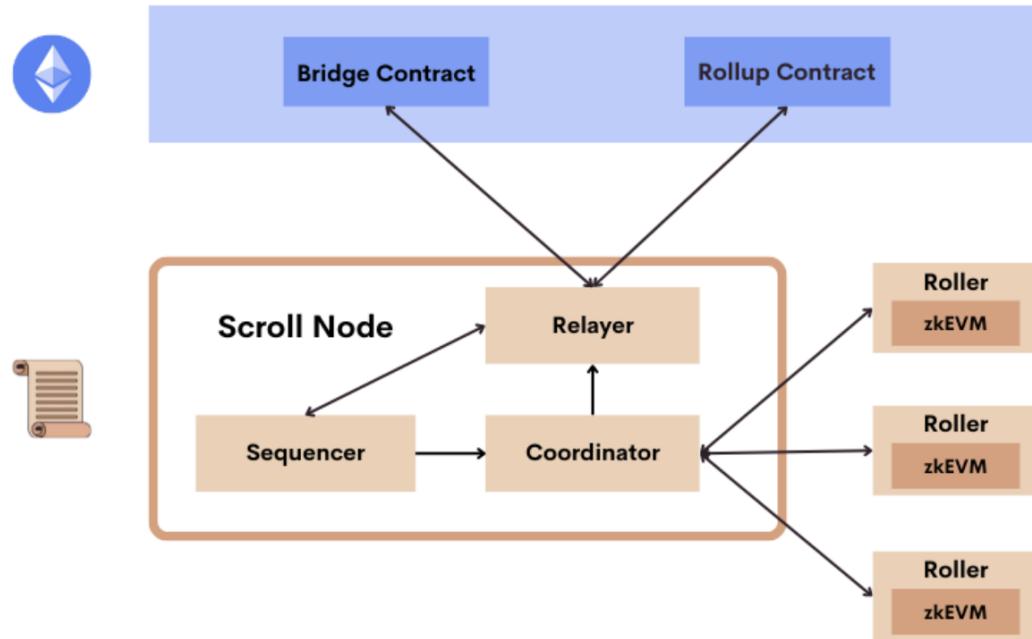
This project is still at an early stage, there alpha testnet will be run on a private PoA fork of Ethereum (the testnet L1) operated by Scroll.

On top of this private chain, will run a testnet Scroll L2 supporting the following features:

- Users will be able to play with a few key demo applications such as a Uniswap fork with familiar web interfaces such as Metamask.
 - Users will be able to view the state of the Scroll testnet via block explorers.
 - Scroll will run a node that supports unlimited read operations (e.g. getting the state of accounts) and user-initiated transactions involving interactions with the pre-deployed demo applications (e.g. transfers of ERC-20 tokens or swaps of tokens).
 - Rollers will generate and aggregate validity proofs for part of the zkEVM circuits to ensure a stable release. In the next testnet phase, we will ramp up this set of zkEVM circuits.
 - Bridging assets between these testnet L1 and L2s will be enabled through a smart contract bridge, though arbitrary message passing will not be supported in this release.
-

From article and article

Scroll Architecture



Components

The **Sequencer** provides a JSON-RPC interface and accepts L2 transactions. Every few seconds, it retrieves a batch of transactions from the L2 mempool and executes them to generate a new L2 block and a new state root.

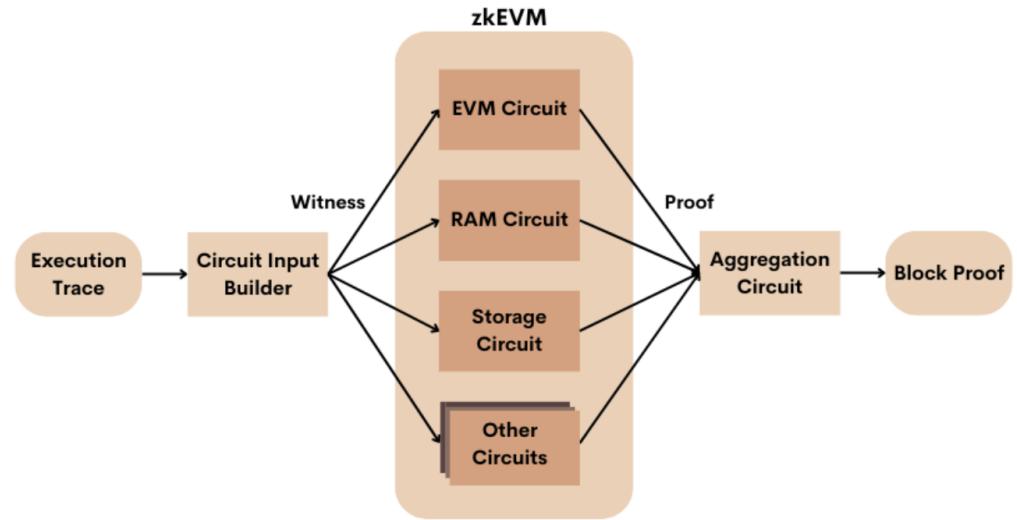
Once a new block is generated, the **Coordinator** is notified and receives the execution trace of this block from the Sequencer.

It then dispatches the execution trace to a randomly-selected **Roller** from the roller pool for proof generation.

The **Relayer** watches the bridge and rollup contracts deployed on both Ethereum and Scroll. It has two main responsibilities.

1. It monitors the rollup contract to keep track of the status of L2 blocks including their data availability and validity proof.
2. It watches the deposit and withdraw events from the bridge contracts deployed on both Ethereum and Scroll and relays the messages from one side to the other.

Rollers - creating proofs



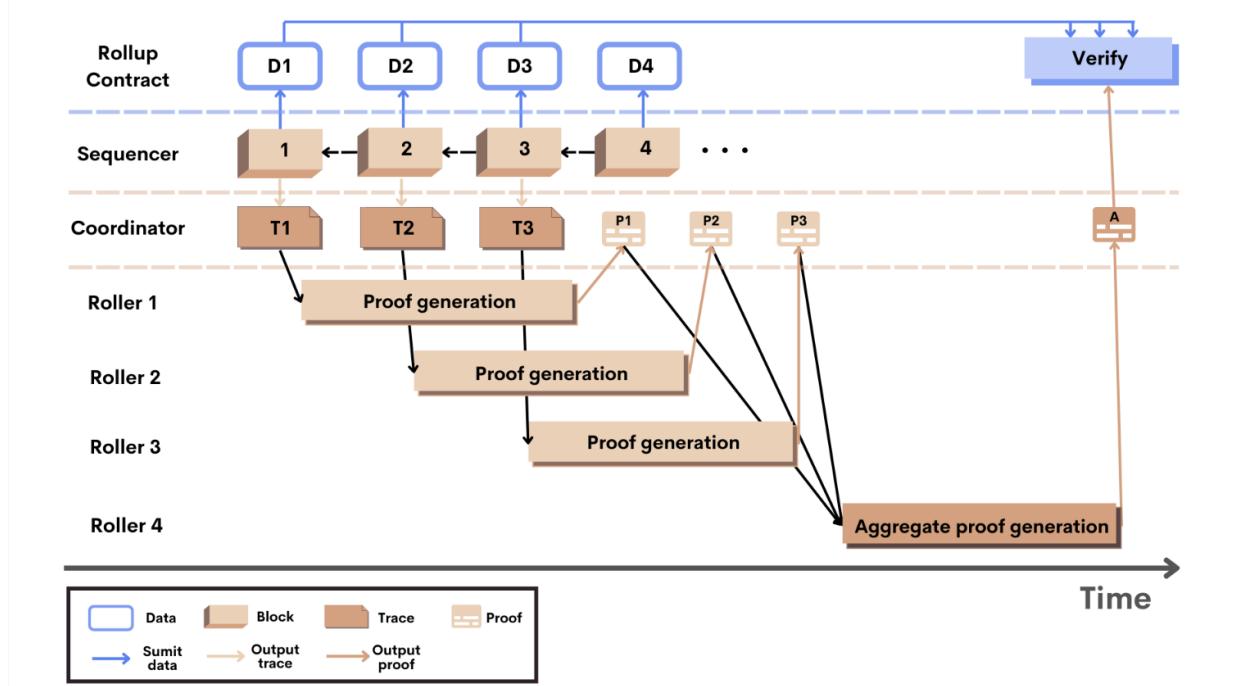
The **Rollers** serve as provers in the network that are responsible for generating validity proofs for the zkRollup

- A Roller first converts the execution trace received from the **Coordinator** to circuit witnesses.
- It generates proofs for each of the **zkEVM** circuits.
- Finally, it uses **proof aggregation** to combine proofs from multiple zkEVM circuits into a single block proof.

The **Rollup contract** on L1 receives L2 state roots and blocks from the Sequencer. It stores state roots in the Ethereum state and L2 block data as Ethereum calldata. This provides **data availability** for Scroll blocks and leverages the security of Ethereum to ensure that indexers including the Scroll Relayer can reconstruct L2 blocks.

Once a block proof establishing the validity of an L2 block has been verified by the Rollup contract, the corresponding block is considered finalized on Scroll.

A useful sequence diagram from the Scroll [Documentation](#)



L2 blocks in Scroll are generated, committed to base layer Ethereum, and finalized in the following sequence of steps:

1. The Sequencer generates a sequence of blocks. For the i -th block, the Sequencer generates an execution trace T and sends it to the Coordinator. Meanwhile, it also submits the transaction data D as calldata to the Rollup contract on Ethereum for data availability and the resulting state roots and commitments to the transaction data to the Rollup contract as state.
2. The Coordinator randomly selects a Roller to generate a validity proof for each block trace. To speed up the proof generation process, proofs for different blocks can be generated in parallel on different Rollers.
3. After generating the block proof P for the i -th block, the Roller sends it back to the Coordinator. Every k blocks, the Coordinator dispatches an aggregation task to another Roller to aggregate k block proofs into a single aggregate proof A .
4. Finally, the Coordinator submits the aggregate proof A to the Rollup contract to finalize L2 blocks $i+1$ to $i+k$ by verifying the aggregate proof against the state roots and transaction data commitments previously submitted to the rollup contract.

Scroll circuit design

1. We need an accumulator to provide the proofs of storage, merkle trees can provide this
2. The execution trace is needed to show the path that the execution took through the bytecode, as this would change because of jumps. This trace is then a witness provided to the circuit.
3. Two proofs are used to show the execution is correct for each opcode
 1. Proof of fetching the data required for the opcode
 2. Proof that the opcode executed correctly.

Scroll are working with Ethereum on this, see this [repo](#) for EVM circuit design, and this design [document](#) from Ethereum.

Design Choices

See this [article](#)

Features

- Community driven development and broad education
 - Be mindful of security and ensure a steady release
 - Decentralisation at all layers is important
 - Develop for Ethereum also
-

Note there are also zk VMs such as [Bonsai](#) from Risc Zero

The RISC Zero zkVM is an open-source, zero-knowledge virtual machine that lets you build trustless, verifiable software in your favourite languages. For more information about our technology, check out the following resources:



The General Purpose Zero-Knowledge VM.
Prove any Computation.
Verify Instantly.



Open Source

Our codebase is licensed under the Apache2 license and includes a full proving and verification system.



Your Favorite Languages

RISC Zero supports Rust and C++ for writing ZK proofs. Any language that compiles to RISC-V can be supported.



Real Microarchitecture

The RISC Zero ZKVM is a verifiable computer that works like a real embedded RISC-V microprocessor, enabling programmers to write ZK proofs like they write any other code.